

Systems Programming

User-Space File System

H. Turgut Uyar Şima Uyar

2009

1 / 21

Topics

User-Space Development

FUSE

Introduction
Read-Only Filesystem
Hello, world

2 / 21

System Programming Levels

- ▶ compiling the kernel:
best performance, every possible functionality
risky, time-consuming
- ▶ kernel modules:
very good performance, less risky, fast development
can not do everything
- ▶ user-space:
even less risky, fast development, can use external libraries
poorer performance, can not do everything

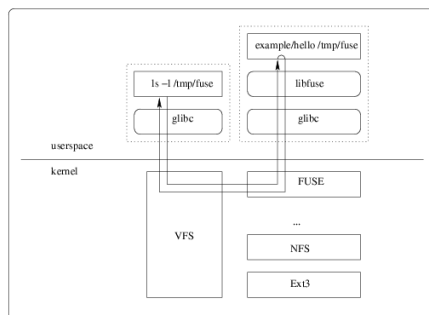
3 / 21

FUSE

- ▶ Filesystem in Userspace
- ▶ develop a file system in user space on top of a kernel module
 - ▶ non-native filesystems (NTFS, ZFS, ...)
 - ▶ changing data storage (SQL, ...)
 - ▶ providing transparent functionality (compression, encryption, ...)
- ▶ two paths for every file/directory:
FUSE path - underlying filesystem path

4 / 21

FUSE Structure



5 / 21

FUSE Development

- ▶ similar to device driver development:
implement system calls
- ▶ needed package: libfuse-dev
- ▶ system calls:
 - ▶ file related:
open, release, read, write, getattr, unlink, ...
 - ▶ directory related:
readdir, mkdir, rmdir, ...

6 / 21

FUSE Development

Example (fuse operations)

```
struct fuse_operations rofs_oper = {
    .getattr = rofs_getattr,
    .readdir = rofs_readdir,
    .mkdir   = rofs_mkdir,
    .unlink  = rofs_unlink,
    .rmdir   = rofs_rmdir,
    .rename  = rofs_rename,
    .open    = rofs_open,
    .read    = rofs_read,
    .write   = rofs_write,
    .release = rofs_release,
    ...
};
```

7 / 21

FUSE Development

Example (path translation)

```
char *rPath = malloc(sizeof(char)*
    (strlen(path) + strlen(rw_path) + 1));

strcpy(rPath, rw_path);
if (rPath[strlen(rPath)-1] == '/') {
    rPath[strlen(rPath)-1] = '\\0';
}
strcat(rPath, path);

return rPath;
```

8 / 21

FUSE Development

directory listing: readdir

```
static int rofs_readdir(
    const char *path,
    void *buf,
    fuse_fill_dir_t filler,
    off_t offset,
    struct fuse_file_info *fi
);
```

9 / 21

FUSE Development

Example (directory listing)

```
upath = translate_path(path);
dp = opendir(upath); /* DIR *dp; */
free(upath);
if (dp == NULL) {
    res = -errno;
    return res;
}

/* fill in the directory info */

closedir(dp);
```

10 / 21

FUSE Development

Example (directory info)

```
/* struct dirent *de; */
while((de = readdir(dp)) != NULL) {
    struct stat st;
    memset(&st, 0, sizeof(st));
    st.st_ino = de->d_ino;
    st.st_mode = de->d_type << 12;
    if (filler(buf, de->d_name, &st, 0))
        break;
}
```

11 / 21

FUSE Development

reading file attributes

```
static int rofs_getattr(
    const char *path,
    struct stat *st_data
);
```

12 / 21

FUSE Development

Example (reading file attributes)

```
upath = translate_path(path);
res = lstat(upath, st_data);
free(upath);
if(res == -1) {
    return -errno;
}
```

13 / 21

FUSE Development

reading from a file

```
static int rofs_read(
    const char *path,
    char *buf,
    size_t size,
    off_t offset,
    struct fuse_file_info *fi)
{
    return 0;
}
```

14 / 21

FUSE Development

Example (reading from a file)

```
upath = translate_path(path);
fd = open(upath, O_RDONLY);
free(upath);
if(fd == -1) {
    res = -errno;
    return res;
}
res = pread(fd, buf, size, offset);
if(res == -1) {
    res = -errno;
}
close(fd);
```

15 / 21

FUSE Development

modification operations

```
static int rofs_mkdir(
    const char *path,
    mode_t mode)
{
    return 0;
}

static int rofs_unlink(const char *path);

/* body */
return -EROFS;
```

16 / 21

FUSE Development

- ▶ compiling:
gcc -o rofs -Wall -ansi -W -std=c99 -g -ggdb
-D_GNU_SOURCE -D_FILE_OFFSET_BITS=64
-lfuse rofs.c
- ▶ mounting:
./rofs <rw_dir> <ro_dir>
- ▶ unmounting:
fusermount -u <ro_dir>
- ▶ running in debug mode:
./rofs -d <rw_dir> <ro_dir>

17 / 21

FUSE Development

Example (fuse operations)

```
static struct fuse_operations hello_oper = {
    .getattr    = hello_getattr,
    .readdir    = hello_readdir,
    .open       = hello_open,
    .read       = hello_read,
};
```

18 / 21

FUSE Development

Example (hello_readdir)

```
if (strcmp(path, "/") != 0)
    return -ENOENT;

filler(buf, ".", NULL, 0);
filler(buf, "..", NULL, 0);
filler(buf, hello_path + 1, NULL, 0);
```

19 / 21

FUSE Development

Example (hello_getattr)

```
memset(stbuf, 0, sizeof(struct stat));
if (strcmp(path, "/") == 0) {
    stbuf->st_mode = S_IFDIR | 0755;
    stbuf->st_nlink = 2;
}
else if (strcmp(path, hello_path) == 0) {
    stbuf->st_mode = S_IFREG | 0444;
    stbuf->st_nlink = 1;
    stbuf->st_size = strlen(hello_str);
}
else
    res = -ENOENT;
```

20 / 21

FUSE Development

Example (hello_read)

```
if (strcmp(path, "/hello") != 0)
    return -ENOENT;

len = strlen("Hello, _world!\n");
if (offset < len) {
    if (offset + size > len)
        size = len - offset;
    memcpy(buf, "Hello, _world!\n"+offset, size);
} else
    size = 0;

return size;
```

21 / 21